

# *Latent Semantic Indexing* application of Ranked Retrieval model

Azza M. N. Abdalghani

Department of Mathematics and Geo-sciences, University of Trieste

*Information Retrieval* course  
Final Project ~ July 2021



# Table of Contents

- 1 LSI : An Overview
- 2 Data Overview
- 3 Implementation Workflow
- 4 Results
- 5 Conclusion

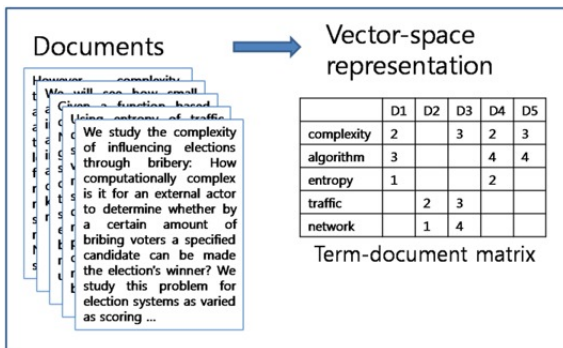
# Latent Semantic Indexing

- Proposed by Deerwester, 1990 .
- **Vector Space** approach for modeling documents .
- Has achieved up to 30% better retrieval performance than lexical searching techniques .
- Advantage : Overcome the problems with term matching retrieval, mainly :
  - **Polysemy** : a term has more than one meaning .
  - **Synonymy** : same topic use different vocabulary.
- Based on a mathematical technique termed **Singular Value Decomposition (SVD)**

- **Documents** and **Queries** are represented as vectors in **t**-dimensional space.
- Map documents and queries into a lower dimensional space : **SVD** and **Dimensionality Reduction**
- Map each **term** into some **concepts**.
- Map **concepts** into **documents**.
- Compute **Similarity** between the documents and queries and return the top-k high ranked documents

# Recall : Vector Space Model

- Represent documents as **vectors** of terms  $d = (t_1, t_2, \dots, t_n)$  where  $t_i (1 \leq i \leq n)$  denoting the single or multiple occurrences of term  $i$  in document  $d \implies$  **Term-Document Matrix** .
- Thus, each term corresponds to a dimension in the space.



# Recall : Singular Value Decomposition

- Goal : **Decompose** the term-document matrix into a product of matrices.

$$M = USV^T$$

- M : original matrix
- U : a term by r matrix
- S : a singular value matrix ( $r \times r$ )
- $V^T$  : a document by r matrix
- where r is the rank of the term-document matrix

# Table of Contents

- 1 LSI : An Overview
- 2 Data Overview
- 3 Implementation Workflow
- 4 Results
- 5 Conclusion

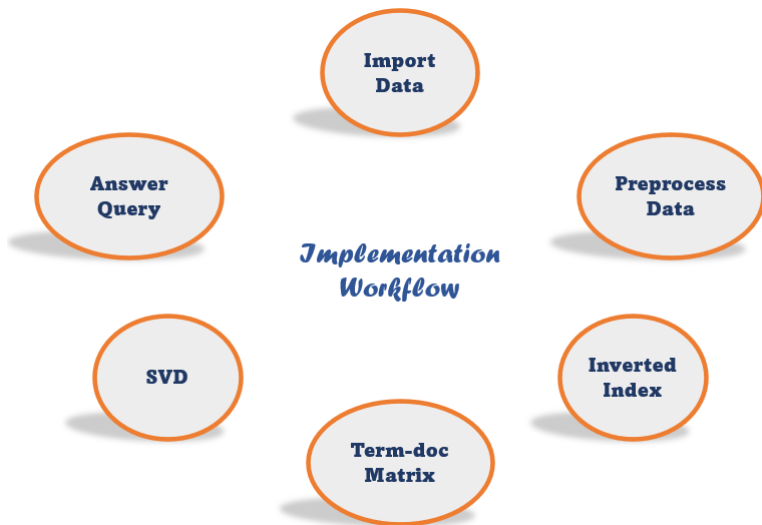
- It is a collection that consists of articles extracted from the magazine Time.
- It includes 5 main files :
  - **README** : Explanation of files
  - **time.all** : the documents
  - **time.que** : the queries
  - **time.rel** : relevance assessments
  - **time.stp** : list of stop words



# Table of Contents

- 1 LSI : An Overview
- 2 Data Overview
- 3 Implementation Workflow**
- 4 Results
- 5 Conclusion

# Implementation Workflow



# Importing and Preprocessing data

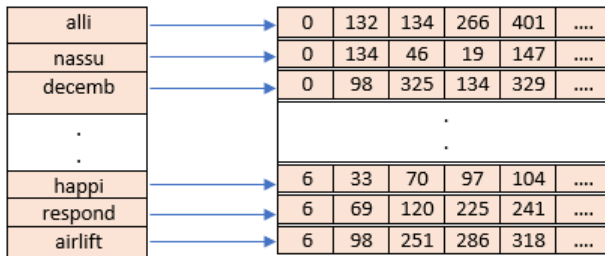
- Import **TIME** dataset : Read TIME.ALL and TIME.STP files.
- Preprocessing data :
  - Tokenization
  - Remove stop-words
  - Remove Punctuation marks
  - Stemming : PorterStemmer()

| Term         | Term after preprocessing |
|--------------|--------------------------|
| EMERGING     | emerg                    |
| CANCELLATION | cancel                   |
| FRANCE       | franc                    |
| DELIVERY     | deliveri                 |
| ANGLOUS      | anglou                   |

- In total, there is 422 articles(documents) and 14952 terms.

# Build Inverted-Index

- Mapping terms to the related documents.
- Build a **dictionary** of terms, connected to a list of **documents ids** that contain the corresponding term.



# Build Term-Document Matrix

- Build a matrix of **M x N**, where rows are terms and columns are documents
- Elements of term-document matrix are **tf-idf** weights

$$TF - IDF_{t,d} = tf_d * idf_t$$

where  $tf_d = \frac{\#term\_occurrences}{|d|}$  and  $idf_t = \log \frac{N}{df_t+1}$

|           | 0        | 1   | ... | 412      | ... | 421      |
|-----------|----------|-----|-----|----------|-----|----------|
| alli      | 0.003114 | 0.0 | ... | 0.001233 | ... | 0.0      |
| nassau    | 0.005462 | 0.0 | ... | 0.0      | ... | 0.0      |
| decemb    | 0.001149 | 0.0 | ... | 0.0      | ... | 0.0      |
| help      | 0.002240 | 0.0 | ... | 0.000887 | ... | 0.0      |
| dictatori | 0.0      | 0.0 | ... | 0.0      | ... | 0.008709 |
| ...       | ...      |     | ... | ...      | ... | ...      |

[14952 rows x 422 columns]

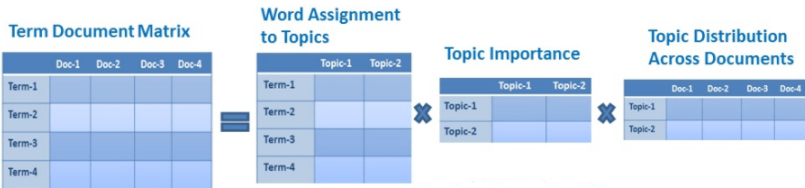
# Matrix Decomposition : SVD

- First, decompose the term-document Matrix into 3 matrices :

$$M = U.\Sigma.V^T$$

- $U$  : Left Singular vectors, called **term-concept matrix**
  - $\Sigma$  : contains Singular values where each value represents a weight of a concept, called **concept matrix**
  - $V^T$  : Right Singular vectors, called **document-concept matrix**
- Then, reduce document dimensions by choosing the number of important concepts **k**.

$$M = U_{M \times k} . \Sigma_{k \times k} . V_{k \times N}^T$$



# How to select the value of $k$ ?

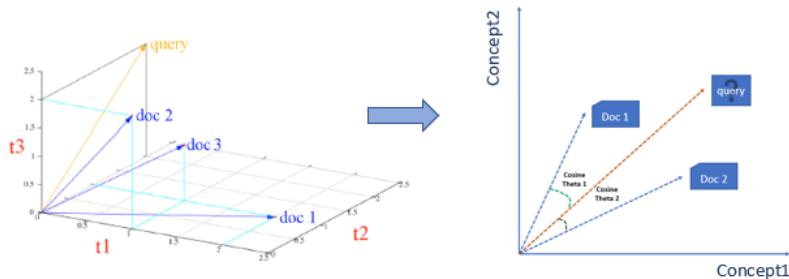
- Determining the number of concepts ( $k$ ) is fully dependent on the corpus.
- In this project, many random values for  $k$  were taken, and generate the results for each one of them.
- Based on the provided queries and relevance assessments, best  $k$  that gives better performance was chosen .

$$k = 215$$

# Answering Query

In order to answer Query :

- Preprocessing query
- Represent query as a vector in terms dimensions
- Remap query in concept space :  $\hat{q} = \Sigma^{-1} U^T q$
- Compute the similarity between each document and the query and assign a score for each document





# How to compute Similarity?

By **Cosine Similarity** : cosine angle between two vectors (query and document), can be computed as the inner product of the two vectors normalized to both have length 1 :

$$sim(v_1, v_2) = \vec{v}_1 \cdot \vec{v}_2$$

Thus, the score of similarity of each document with query can be :

$$score(q, d_j) = \vec{q} \cdot \vec{d}_j$$

Finally, Scores get sorted , higher score means more relevant document (collinear with query, e.i point to the same direction), return the top k ranked documents.

# Table of Contents

- 1 LSI : An Overview
- 2 Data Overview
- 3 Implementation Workflow
- 4 Results**
- 5 Conclusion

- For testing our implementation, a set of queries were taken with their corresponding answers from **time.que** and **time.rel**.
- A list of relevant documents and their scores were retrieved as an answer to each query.
- Finally, as a performance metric, **R-Precision** was computed for each query and at the end compute the **Average R-Precision (ARP)** for the set of queries.

$$\implies ARP = \sum_n RP_n$$

where n is number of queries

# Experimental Results

```
#69 70 100 115 121 139 159 194 210 224 234 309 379 388
```

```
query_1 = ["THE BAATH (RENAISSANCE) PARTY FOUNDED BY MICHEL AFLAK, WHICH HAS GAINED CONTROL OF SYRIA AND IRAQ AND AIMS TO UNITE  
M_decomposition = LSI(M, 215)  
relevant_docs_1 = answer_query(query_1,M_decomposition,terms,13)  
print(relevant_docs_1)
```

```
[[377, 0.9074748457722953], [386, 0.7648667535593023], [309, 0.5648373281418607], [194, 0.5316742024010941], [115, 0.4927082195  
7710237], [159, 0.4702934871295355], [210, 0.3849730137900559], [100, 0.35748794415113494], [121, 0.3377839712475548], [70, 0.3  
3242266716334884], [234, 0.326790886761082], [139, 0.291839556073581], [224, 0.21395117397533003]]
```

It is observed that 11 of the retrieved documents are relevant

$$\rightarrow R - Precision_{q_1} = \frac{11}{13} = 0.84615$$

```
#12 61 155 156 242 269 339 358
```

```
query_6 = ["CONTROVERSY BETWEEN INDONESIA AND MALAYA ON THE PROPOSED FEDERATION OF MALAYSIA, WHICH WOULD UNITE FIVE TERRITORIES  
M_decomposition = LSI(M, 215)  
relevant_docs_6 = answer_query(query_6,M_decomposition,terms,7)  
print(relevant_docs_6)
```

```
[[61, 0.6785103101149405], [358, 0.6234520083076952], [156, 0.60534689432005], [155, 0.5143008007463842], [303, 0.4370222262841  
8297], [269, 0.36305323227580283], [339, 0.3591404003504791]]
```

It is observed that 6 of the retrieved documents are relevant

$$\rightarrow R - Precision_{q_6} = \frac{6}{7} = 0.85714$$

# Experimental Results

```
#39 22 73 173 189 219 265 277 360 396
query_3 = ["COALITION GOVERNMENT TO BE FORMED IN ITALY BY THE LEFT-WING SOCIALISTS, THE REPUBLICANS, SOCIAL DEMOCRATS, AND CHRIS'
M_decomposition = LSI(M, 215)
relevant_docs_3 = answer_query(query_3,M_decomposition,terms,8)
print(relevant_docs_3)
```

[[277, 0.7656766098060409], [360, 0.7491591189751454], [394, 0.6779197006972426], [265, 0.5678370396247605], [219, 0.4940555543654217], [189, 0.44341632087842253], [134, 0.40568774969918425], [22, 0.39216211767542447]]

$$\rightarrow R - Precision_{q_3} = \frac{6}{8} = 0.75$$

```
#27 272 295 306
query_8 = ["BRITISH PROPOSAL FOR NEW HIGH LEVEL NEGOTIATIONS WITH RUSSIA OR A FOUR-POWER SUMMIT MEETING ."]
M_decomposition = LSI(M, 215)
relevant_docs_8 = answer_query(query_8,M_decomposition,terms,3)
print(relevant_docs_8)
```

[[151, 0.4218138316235267], [111, 0.41488162259210043], [306, 0.31681114968467844]]

$$\rightarrow R - Precision_{q_8} = \frac{1}{3} = 0.33333$$

For the 8 queries, we obtain  $ARP = 0.64247$

# Table of Contents

- 1 LSI : An Overview
- 2 Data Overview
- 3 Implementation Workflow
- 4 Results
- 5 Conclusion**

- LSI attempts to overcome the common problems of search engines (Synonymy and Polysemy)
- LSI designed to uncover the latent semantic structure of a document space by building a semantic space.
- Documents and queries are represented in this semantic space and a similarity score is measured.
- **Drawbacks:**
  - SVD is expensive to compute
  - Requires re-computing SVD when new documents arrive



**Thank you for your  
attention 😊**